

Graph Neural Network Adapters for Multi-Label Chest X-Ray Classification Without VLM Fine-Tuning: Frozen Scorers, Structured Heads, and Leakage-Free Calibration

Repository: MBZAI multi-label CXR pipeline **Codebase:** scripts/01–13, configs/, gradio_inference.py **Dataset:** CheXpert-v1.0-small (Frontal + Lateral) **Label space:** {Atelectasis, Cardiomegaly, Effusion, Pneumonia, Edema, Consolidation, No Finding} ($C = 7$)

Positioning. The central proposal is **domain adaptation for multi-label classification without fine-tuning the VLM**: keep the foundation model frozen (one-time inference or API scores), and learn a **small graph-based adapter** that maps VLM outputs—and optionally a second frozen image encoder—into calibrated logits for the target domain (here, CheXpert-style labels and masking). That trades **full VLM fine-tuning** (GPU memory, catastrophic forgetting, data-hungry updates) for **adapter training** on cheap tabular tensors $(x_{\text{logits}}, x_{\text{probs}}, y_{\text{true}}, y_{\text{mask}})$ plus optional CLIP embeddings.

Abstract

We **propose graph neural network (GNN) adapters as a practical alternative to end-to-end VLM fine-tuning** for multi-label chest X-ray classification: the VLM stays frozen while a lightweight head performs **domain-specific calibration and structured reasoning** over labels. Concretely, we re-calibrate a frozen Qwen2-VL-2B-Instruct baseline on CheXpert and compare four adapters of increasing structure: (i) an MLP over VLM logit/probability vectors; (ii) a residual label-graph GNN on a *co-error* adjacency mined from training disagreements; (iii) a homogeneous CLIP-conditioned label-graph GNN; and (iv) a bipartite *attribute* $\rightarrow$ *object* GNN with frozen CLIP object features and VLM attribute nodes. All adapters use masked, class-weighted binary cross-entropy on patient-grouped splits and are evaluated under (a) fixed $t = 0.5$ and (b) a leakage-free 4-way *train/calib/val/test* protocol with per-class thresholds tuned **only** on calib. The bipartite CLIP variant (gnn13_clip_bipartite) reaches the best calibrated test macro-F1 of **0.6889**, **+3.4** over the calibrated MLP (**0.6544**) and **+3.8** over the same leakage-free calibrated reference applied to frozen VLM probabilities (**0.6512**; i.e. exported x_{probs} with thresholds fit on calib only—see §6.1). Naive $t = 0.5$ macro-F1 on the frozen model remains ≈ 0.047 because logits are systematically negative off No Finding. We also show how threshold tuning on the same split as evaluation can inflate macro-F1 for poorly calibrated logits—a methodological artifact, not model capacity—and document how to avoid it with a held-out calib split (§6.2).

1. Introduction

Thesis. Fine-tuning a large VLM for every new hospital, label set, or policy change is often **prohibitively expensive**: it requires gradient updates through billions of

parameters, careful replay or regularization to limit forgetting, and substantial curated data. A complementary strategy is **post-hoc domain adaptation at the decision layer**: treat the VLM as a fixed feature generator $f_{\text{VLM}}(x) \mapsto (z, p) \in \mathbb{R}^{2C}$ (and optionally a second frozen encoder $e_{\text{CLIP}}(I)$), then learn a small module g_θ with $|\theta| \ll |\text{VLM}|$ so that $\hat{z} = g_\theta(z, p, e)$ matches the target domain’s labels and masking. **GNN-based** g_θ is attractive because multi-label CXR findings are **not independent**: co-occurrence, mutual exclusion, and systematic error patterns are naturally expressed as **message passing on a label graph** or as **attribute→image bipartite** structure—inductive biases that a flat MLP must rediscover from data alone.

Multi-label classification of chest X-rays under the CheXpert label policy still presents three coupled difficulties even under this adapter framing: (1) extreme **class imbalance**, (2) **uncertain (-1) labels** that must be policy-mapped, and (3) **miscalibrated** zero-shot VLM probabilities (often dominated by No Finding at naive thresholds). We keep the VLM frozen and ask three empirical questions that stress-test the **GNN-as-domain-adapter** idea:

1. **Is structure useful?** Does encoding label co-occurrence/co-error as a graph help over a flat MLP that already sees the full $[x_{\text{logit}}; x_{\text{prob}}]$ vector?
2. **Is the image useful again?** Once we already have VLM scores, is there marginal value in re-injecting a *different* frozen image encoder (CLIP)?
3. **Is the wiring useful?** Does the bipartite *attribute* → *object* topology provide a better inductive bias than a homogeneous label graph?

To answer these in a way that is **not contaminated by threshold-tuning leakage**, we build a 4-way patient-grouped split (train_fit / calib / val / test) and frozen per-class thresholds that are *only* tuned on calib and then re-applied unchanged to val and test. All five models in our comparison (zero-shot, MLP, three GNNs) are evaluated under this identical protocol.

1.1 What “domain adaptation without finetuning” buys you

Aspect	Full VLM fine-tuning	GNN (or MLP) adapter on frozen (z, p)
Trainable parameters	10^9 - 10^{10}	10^4 - 10^6 in this repo
GPU memory at train	Very high (vision + LM)	Low (small tensors + optional CLIP cache)
Forgetting / drift	Risk when updating foundation	None on the VLM; only θ changes
New label policy or site	Often re-finetune	Re-align JSONL + retrain adapter + retune thresholds
Domain knowledge	Implicit in gradients	Explicit in graph topology (co-error, bipartite)

The **GNN variants** add one more lever: **structured transfer**—sharing statistical strength across labels via edges—without ever backpropagating into the VLM. That

is the sense in which this work is a **proposal for cheap domain adaptation** rather than a claim that GNNs always beat MLPs (under fair calibration, residual label-graph and MLP are close; CLIP + bipartite is where structure pays).

Contributions

- **Conceptual proposal: GNN-based adapters as a frozen-VLM alternative to fine-tuning** for multi-label CXR: adapt the *output distribution and label dependencies* to the target domain while the foundation model remains a fixed scorer.
- **Adapter family** spanning four distinct inductive biases over the same frozen VLM outputs (vlm_zeroshot, vlm_mlp, gnn07_label_residual, gnn12_clip_vlm_homo, gnn13_clip_bipartite), each registered in a small model registry (scripts/model_registry.py) for organized, reproducible run management.
- **Co-error label graph**: a sparsified, row-normalized directed adjacency built from VLM disagreements on the training split (scripts/04_build_coerror_graph.py), which serves as the structural prior for the residual and homogeneous GNNs.
- **Bipartite NativeGNN**: a CLIP-object / VLM-attribute message-passing network with optional VLM-positive edge masking (scripts/gnn_bipartite.py, scripts/13_train_bipartite_gnn_adapter.py).
- **Fairness-aware evaluation**: a 4-way calibrated protocol (scripts/03_make_multilabel_split, scripts/08_tune_thresholds.py, scripts/09_evaluate_test.py) that quantifies and removes the threshold-tuning leakage that otherwise inflates GNN scores by tens of F1 points.
- **Empirical findings** reported in §6, with a root-cause analysis (§7) of the most common evaluation pitfalls and the genuine gains of CLIP injection and bipartite wiring.

2. Related Work

CheXpert and uncertain labels. CheXpert (Irvin et al., 2019) introduced the U-Ones / U-Zeros / U-Ignore policies for the -1 (uncertain) label. We use **U-Zeros** with explicit masking: empty entries set both y and $mask$ to 0, so excluded samples contribute nothing to the loss or to per-class F1.

Frozen-foundation adapters. A growing body of work freezes a large pretrained encoder and trains a small adapter for downstream tasks (LoRA, prefix tuning, linear probes). For multi-label classification specifically, Chen et al. (2019, ML-GCN) proposed using a label-correlation graph to refine class-wise predictions. Our gnn07 variant is a streamlined residual realization of this idea: nodes are labels, features are per-row $(\text{logit}_i, \text{prob}_i)$, and a single message-passing step adds a residual on top of the VLM logits.

CLIP for medical imaging. CLIP (Radford et al., 2021) is not a chest-X-ray-specialised encoder, but its image embeddings are a cheap, complementary signal to those of any other VLM. We use a frozen openai/clip-vit-base-patch32 image

branch and project it into the GNN node space (gnn12) or into a dedicated *object* node (gnn13).

Bipartite attribute graphs. Encoding *instances* and *targets* as different node roles is classic in heterogeneous and relational graph networks: message passing can traverse **bipartite incidence structure** rather than enforcing a homogeneous $C \times C$ label graph everywhere. Foundations include **relational GCN** for typed edges between entity kinds (Schlichtkrull et al., 2018), **semantic-level meta-path attentive aggregation** across multiple node roles (Wang et al., 2019, HAN), and the modular **Neural Message Passing** view in which pairwise updates generalize heterogeneous interaction patterns (Gilmer et al., 2017). Our bipartite instantiation places **one CLIP-derived object node per image** and C **VLM-attribute nodes**, with weighted-mean attribute→object pooling before readout—a lightweight inductive bias aligned with CXR semantics without depending on PyG (scripts/gnn_bipartite.py).

Threshold calibration. Per-label operating points tie directly to probabilistic calibration and **proper scoring** viewpoints (Niculescu-Mizil & Caruana, 2005; Guo et al., 2017). Selecting thresholds by grid search remains common in multi-label setups (Lipton et al., 2014 optimise F-measure via threshold choices). Critically: **fitting thresholds—or any discrete decision rule—on the same split later used as a “selection” metric** induces optimistic bias (Varma & Simon, 2006; Cawley & Talbot, 2010). Our **calib** split mirrors the textbook fix used in unbiased model comparison: isolate post-hoc threshold fitting on disjoint data, unchanged when applied to validation and held-out test. §6 quantifies how badly that discipline matters when logits are saturated.

3. Dataset and Preprocessing

3.1 Source

We use **CheXpert-v1.0-small** with the standard `train.csv` (≈ 223 k rows after header) and `valid.csv` (234 rows). Both are CSVs with one row per study + view; each row contains 14 finding columns. We restrict to the 7-label canonical space below.

3.2 Canonical label mapping (scripts/common_multilabel.py)

```
"Atelectasis"      -> "Atelectasis"
"Cardiomegaly"     -> "Cardiomegaly"
"Pleural Effusion" -> "Effusion"
"Pneumonia"        -> "Pneumonia"
"Edema"            -> "Edema"
"Consolidation"    -> "Consolidation"
"No Finding"       -> "No Finding"
```

Uncertain (-1) labels are mapped via `parse_uncertain(value, policy="u_zeros")`: $1 \rightarrow (1,1)$, $0 \rightarrow (0,1)$, $-1 \rightarrow (0,1)$ under U-Zeros, and `empty` $\rightarrow (0,0)$. The `mask=0` entries are excluded from both loss and F1 computation, which is the single most important detail for honest multi-label evaluation on CheXpert.

3.3 VLM alignment (scripts/02_align_vlm_outputs.py)

We score every image once with **Qwen2-VL-2B-Instruct** (zero-shot, structured-prompt-and-parse) and persist {path, scores: {label: prob}} to JSONL shards under data/outputs_vlm_corrected/. Alignment joins on normalize_path(path) and produces aligned rows of the form

```
{path, image_id, patient_id,
 x_probs   in R^7,      % in [0,1]
 x_logits  in R^7,      % safe_logit(x_probs), eps 1e-6
 y_true    in {0,1}^7,
 y_mask    in {0,1}^7}
```

This (x_logits, x_probs, y_true, y_mask) quadruple is the only thing the four downstream adapters see during training; the VLM is never re-invoked at adapter-train time.

3.4 Splits

Two split protocols co-exist in the repo, sharing the same row schema and the same patient_id grouping (no patient overlap across splits).

(a) 3-way default (scripts/03_make_multilabel_splits.py): 70 / 15 / 15 patient-level shuffle, seed 42. Used for all protocol=default numbers.

(b) 4-way calibrated4way (scripts/03_make_multilabel_splits_4way.py): 70 / 10 / 10 / 10 patient-level shuffle, seed 42 → train_fit / calib / val / test. The calib split exists *only* to host per-class threshold tuning; val and test are entirely unseen by the threshold optimizer.

Multi-label intermediates produced by these scripts carry on the order of **44k train** rows under the patient-grouped regimes used here (exact counts drift slightly with 02_align_vlm_outputs.py parity); registry runs report **43 778 train / 9 357 val / 9 197 test** for protocol=default and analogous 4-way sizes for calibrated4way (see metrics.json dataset_sizes under each run).

3.5 Co-error label graph (scripts/04_build_coerror_graph.py)

For every training row with at least one positive label, we form

$$\text{present}(r) = \{i : y_i = 1 \wedge m_i = 1\}, \quad \text{absent}(r) = \{j : y_j = 0 \wedge m_j = 1\},$$

and accumulate

```
M[i, j] += 1   for i in present sorted ASC by VLM prob (top_k=3)
                for j in absent  sorted DESC by VLM prob (top_k=3), j != i
```

i.e. an edge $i \rightarrow j$ is added whenever the VLM **missed** a true positive i while **hallucinating** a likely false positive j . Rows of M are L1-normalized to give W , then sparsified per-source to the top- $k=3$ targets with weight $\geq \tau=0.02$. The resulting (edge_index, edge_weight) defines an asymmetric directed graph stored in

data/processed/graph/{edge_index, edge_weight}.json. The same graph is reused, identically, by gnn07 and gnn12.

4. Methodology

4.1 Problem and notation

For a single image with VLM outputs $z = \text{logits} \in \mathbb{R}^C$, $p = \sigma(z) \in [0, 1]^C$, label vector $y \in \{0, 1\}^C$ and mask $m \in \{0, 1\}^C$, each adapter outputs **calibrated logits** $\hat{z} \in \mathbb{R}^C$ and is trained with masked, class-weighted BCE:

$$\mathcal{L}(\hat{z}, y, m) = \frac{\sum_{i=1}^C m_i \cdot \text{BCEWithLogits}(\hat{z}_i, y_i; w_i^+)}{\sum_{i=1}^C m_i + \varepsilon}, \quad w_i^+ = \min\left(\frac{N_i^-}{\max(N_i^+, 1)}, 100\right).$$

N_i^+, N_i^- are training positive/negative counts on the (masked) train split; w^+ is the pos_weight used by binary_cross_entropy_with_logits.

4.2 M0 — vlm_zeroshot (VLMZeroShot)

$$\hat{z} = z, \quad \hat{p} = \sigma(z)$$

No learnable parameters. Decision rule: $\hat{y}_i = \mathbf{1}[\hat{p}_i \geq 0.5]$. Implemented in scripts/05_run_baseline_frozen_vlm.py.

4.3 M1 — vlm_mlp (VLMFeatureMLP)

A flat 2-layer MLP over the concatenated logit/prob vector:

$$x = [z_1, p_1, z_2, p_2, \dots, z_C, p_C] \in \mathbb{R}^{2C}, \quad \hat{z} = W_2 \text{Dropout}(\text{ReLU}(W_1 x)),$$

with $W_1 \in \mathbb{R}^{64 \times 2C}$, $W_2 \in \mathbb{R}^{C \times 64}$, AdamW, lr=1e-3, weight_decay=1e-4, 20 epochs, dropout 0.1. Implemented in scripts/06_run_baseline_mlp.py. This is the simplest baseline that can learn class-specific *bias correction* and per-label *temperature*.

4.4 M2 — gnn07_label_residual (LabelGraphResidualGNN)

C -node homogeneous graph with row-normalized adjacency $A \in \mathbb{R}^{C \times C}$ (built in build_adj by adding self-loops then row-stochastic normalization). Per-row node features are 2-d: $h_i^{(0)} = [z_i, p_i]$. The model is a tiny per-node MLP followed by a single message-pass and a residual on the original logits:

$$g_i = W_2 \text{ReLU}(W_1 h_i^{(0)}) \in \mathbb{R}, \quad \Delta = g A^\top \in \mathbb{R}^C, \quad \hat{z} = z + \alpha \Delta.$$

hidden_dim=32, alpha=0.5, lr=3e-4, AdamW with cosine LR + 2-epoch warmup, gradient clip 1.0, 80 epochs max with early_stop_patience=18 on val_bce. Implemented in scripts/07_train_gnn_adapter.py.

4.5 M3 — gnn12_clip_vlm_homo (ClipVlmHomogeneousGNN)

Same homogeneous label graph as M2, but each node now also sees a projected CLIP image embedding. With $e \in \mathbb{R}^{D_{\text{clip}}}$ the frozen CLIP image embedding (openai/clip-vit-base-patch32):

$$\tilde{z} = \text{ReLU}(W_e e) \in \mathbb{R}^H, \quad h_i^{(0)} = \text{ReLU}(W_n[\tilde{z}; z_i; p_i]).$$

Then K GNN layers with normalized adjacency A (elements A_{ij}):

$$h_i^{(k+1)} = \text{ReLU}\left(W^{(k)} \sum_j A_{ij} h_j^{(k)}\right), \quad \Delta_i = W_h h_i^{(K)} \in \mathbb{R}, \quad \hat{z}_i = z_i + \alpha \Delta_i \quad (i = 1, \dots, C).$$

Writing $\Delta = (\Delta_1, \dots, \Delta_C)^\top$ yields $\hat{z} = z + \alpha \Delta \in \mathbb{R}^C$ componentwise.

hidden_dim=64, gnn_layers=2, alpha=0.5, lr=3e-4, AdamW + cosine, batch 32, 60 epochs, early-stop 16. Implemented in scripts/12_train_clip_vlm_gnn_adapter.py. CLIP embeddings are precomputed once per split and cached as a single .pt (--clip_cache_pt).

4.6 M4 — gnn13_clip_bipartite (ClipBipartiteAttributeGNN)

A **bipartite** graph with C attribute nodes (one per label) and one object node (the image). Attribute features are still $[z_i, p_i]$; the object feature is a linear projection of CLIP: $o^{(0)} = W_{\text{clip}} e$. Each bipartite layer aggregates a weighted-mean message from attributes to the object, projects, concatenates with the object state, then updates with a small MLP + dropout (scripts/gnn_bipartite.py::BipartiteMessagePassingLayer):

$$\mu = \frac{\sum_i w_i W_{\text{am}} a_i}{\sum_i w_i + \varepsilon}, \quad \nu = W_{\text{ap}} \mu, \quad o^{(k+1)} = \text{Dropout}(\text{ReLU}(W_u[o^{(k)}; \nu])).$$

After L layers (hidden_dims=[512,256]), a single classifier head produces C logits from the object state, and a residual term re-anchors them to the VLM:

$$\hat{z} = W_{\text{cls}} o^{(L)} + \alpha z, \quad \alpha = 0.5.$$

The edge weights w_i come from build_bipartite_edge_weights(p, mode, τ):

- mode=all: $w_i = 1$ (uniform mean).
- mode=vlm_positive: $w_i = \mathbf{1}[p_i \geq \tau]$, with all-ones fallback if a row has no edges.

object_feature_dim=512, dropout 0.2, lr=3e-4, AdamW + cosine, batch 32, 60 epochs, early-stop 16. Implemented in scripts/13_train_bipartite_gnn_adapter.py.

4.7 Why these four shapes

Adapter	Sees image again?	Uses label graph?	Uses CLIP?	Topology
vlm_mlp	no	no	no	dense flat
gnn07_label_resolution	yes	yes ($A_{C \times C}$)	no	homogeneous
gnn12_clip_vlm_hypos	yes	yes ($A_{C \times C}$)	yes	homogeneous, image-broadcast
gnn13_clip_bipartite	yes	implicit via attr→obj	yes	bipartite

Each row adds **exactly one** capability over the previous one, which lets §6 attribute the lift to a single change at a time.

5. Experimental Setup

5.1 Hardware and software

CUDA-only training. PyTorch 2.x with `torch.cuda` enforced (scripts/{06,07,12,13}*.py raise if CUDA is unavailable). HuggingFace transformers for CLIP vision encoding. Runs are GPU-pinned via `--gpu_id`.

5.2 Training protocols

Every training script supports the same pair of CLI knobs:

```
--protocol {default | calibrated4way}
--run_id    <unique tag>                # auto-generated if omitted
--resume_from <checkpoint.pt>          # warm start
```

Outputs are written to `data/processed/experiments/<model_id>/<protocol>/<run_id>/`, with `runs_index.json`, `latest.json`, `best.json` pointers maintained by `scripts/model_registry.py`.

5.3 Evaluation protocols

We always report **masked macro-F1** (per-class F1 on rows where `mask=1`, then averaged across the 7 classes — `common_multilabel.f1_from_counts` and `scripts/06,07,12,13 masked_macro_f1`).

Two thresholding modes are reported side-by-side:

- **@0.5:** $\hat{y}_i = \mathbf{1}[\hat{p}_i \geq 0.5]$ for every class.
- **@per_class_thr:** per-class thresholds picked by `scripts/08_tune_thresholds.py` via a grid sweep $t \in \{0.05, 0.10, \dots, 0.95\}$ that maximises class F1 on a *calibration* prediction set.

The leakage-free calibrated4way protocol is the recommended one and is the protocol used to declare the *best* model in this report. It enforces:

1. tune `per_class_thresholds.json` **only** on `calib_predictions.json`,
2. apply those frozen thresholds to `val_predictions.json` and `test_predictions.json` via `scripts/09_evaluate_test.py`,
3. compare models on the resulting `test_metrics_calibrated.json`.

Conceptually this is **hold-out threshold calibration** analogous to reserving part of labeled data purely for deploying a decision rule—a standard guard against inflated metrics when optimisation and reporting coincide (Varma & Simon, 2006; Cawley & Talbot, 2010; full entries in Appendix C).

5.4 Hyperparameters and selection

Best-checkpoint selection is `--best_metric val_bce` everywhere by default. We deliberately do **not** select on `val_macro_f1@thr` because (a) @0.5 macro-F1 is degenerate for a miscalibrated VLM and (b) selecting on `@thr` re-introduces a soft form of threshold-tuning leakage. Cosine LR with 2-epoch linear warmup, `min_lr=1e-6`, gradient-norm clip 1.0, AdamW with `weight_decay=1e-4`, seed 42 throughout.

5.5 Reproducibility

End-to-end pipeline:

```
# One-shot (recommended): splits, graphs, dual CLIP caches, all models, calibrated + le
./scripts/reproduce_all_results.sh
# Artifact layout: RUN_ID=<tag> appears under data/processed/experiments/<model_id>/<pr
```

Or step-by-step (must use **separate** CLIP caches for default vs calibrated4way rows so path order matches the rows JSON):

```
python scripts/01_build_canonical_labels.py
python scripts/02_align_vlm_outputs.py
python scripts/03_make_multilabel_splits.py
python scripts/03_make_multilabel_splits_4way.py
python scripts/04_build_coerror_graph.py --train_rows_json data/processed/splits/train_
python scripts/04_build_coerror_graph.py --train_rows_json data/processed/splits_4way/t
python scripts/05_run_baseline_frozen_vlm.py ...
python scripts/06_run_baseline_mlp.py ...
python scripts/07_train_gnn_adapter.py ...
python scripts/12_train_clip_vlm_gnn_adapter.py --clip_cache_pt data/processed/embeddin
python scripts/12_train_clip_vlm_gnn_adapter.py --clip_cache_pt data/processed/embeddin
python scripts/13_train_bipartite_gnn_adapter.py --clip_cache_pt ... # mirror 12 - def
python scripts/08_tune_thresholds.py ...
python scripts/09_evaluate_test.py ...
python scripts/11_package_report.py
```

Canonical comparison numbers in §6 mirror **reports/comparison/overall.json** produced by **python scripts/11_package_report.py**, aggregated from **best.json** → `.../<model_id>/<protocol>/repro_full_20260503/` (tables round to four decimals).

Run inventory: data/processed/experiments/<model_id>/<protocol>/{runs_index,latest,best}
 Reports: reports/comparison/overall.md, reports/comparison/overall.json, reports/gnn_adapter/report.md. Legacy checkpoints and superseded caches may live under data/processed/experiments/_archive_20260503/ if you archived older runs.

6. Results

All numbers below are macro-F1, rounded to 4 decimals; raw 6-decimal values are in runs_index.json and the metrics.json of each run directory. The default columns use the 3-way split; the calibrated4way column uses the leakage-free 4-way split with thresholds tuned on calib.

6.1 Main comparison

Model	Default Val @0.5	Default Test @0.5	Default Val @thr	Default Test @thr	Calib4way Val	Calib4way Test
vlm_zeroshot (frozen VLM)	0.0473	0.0472	NA	NA	0.6513	0.6512*
vlm_mlp (MLP adapter)	0.5208	0.5191	NA	NA	0.6548	0.6544
gnn07_label_residual	0.0442	0.0423	0.0442	0.0423	0.6513	0.6512
gnn12_clip_vlm	0.6095	0.6013	0.6095	0.6013	0.6792	0.6777
gnn13_clip_bipartite	0.6542	0.6371	0.6542	0.6371	0.6923	0.6889

* **Calib4way** column: macro-F1 after 08_tune_thresholds.py on calib_predictions.json only, then evaluated on **val/test** predictions with those frozen thresholds (test_metrics_calibrated). For **vlm_zeroshot**, predictions are the raw frozen VLM probabilities (x_probs) exported via scripts/export_rows_to_predictions.py; this is **not** the same numeric story as $\hat{y} = \mathbb{1}[p \geq 0.5]$, where macro-F1 is ≈ 0.047 (near-trivial negatives). Detailed @0.5 JSON for the auxiliary baseline path is still in data/processed/experiments/baseline_frozen_

Headline result. Under the leakage-free calibrated protocol (reports/comparison/overall.json), gnn13_clip_bipartite reaches **macro-F1 = 0.6889** on test (+3.4 macro-F1 **points** vs MLP **0.6544**, +1.1 vs gnn12 **0.6777**). Calibrated zeroshot, residual GNN, and **MLP cluster near 0.6512**, leaving bipartite +0.0377 calibrated macro-F1 over that reference (**0.6512** $\rightarrow$ **0.6889**). On the **default** splits, masking at $t = 0.5$ leaves frozen zeroshot at ≈ 0.047 macro-F1, while trained heads move into the ≈ 0.5 -0.65 @0.5 band (same table).

6.2 The threshold-tuning leakage trap (RCA)

Look at gnn07_label_residual on the replicated **default** run (repro_full_20260503): at fixed threshold 0.5, masked macro-F1 on **validation** stays near **0.044** (metrics.json).

If thresholds are tuned on `val_predictions.json` and evaluated **again on validation** (`.../val_metrics_thr_tuned_on_val_LEAKY.json`), macro-F1 on that same split climbs to ≈ 0.657 — an $\approx +0.61$ artefact from re-using the same labelled split both to pick thresholds and to report the headline number. Mechanism:

1. The residual adapter learns a near-zero-mean correction Δ on top of frozen VLM logits whose distribution is roughly $\mathcal{N}(\mu \ll 0, \sigma)$ for non-No Finding classes (Qwen2-VL’s soft *negative* bias).
2. Sigmoid of those logits hugs ≈ 0.05 – 0.20 for true positives. At $t = 0.5$ virtually no class fires $\rightarrow$ recall $\approx 0 \rightarrow$ F1 ≈ 0 .
3. Threshold sweep $t \in \{0.05, \dots, 0.95\}$ recovers each class’s F1 by simply picking a low threshold; tuning *and* reporting on the same split optimistically samples the F1-maximising operating point.

Compare this to the leakage-free 4-way protocol on the same model (`test_metrics_calibrated.json`) **0.6512** on test in the replicated registry—that is the leakage-free calibrated score, and **it matches calibrated frozen probabilities (0.6512)** in this artifact bundle. The qualitative lesson persists: Δ logits are saturated at $t = 0.5$, whereas per-threshold optimisation on p restores recall; tuning those thresholds **on evaluation data** exaggerates perceived lift. Always tune thresholds on a disjoint calib split reserved for scoring rules alone.

6.3 Effect of CLIP image features (M2 $\rightarrow$ M3)

Holding the homogeneous label graph fixed and adding a frozen CLIP branch lifts calibrated test F1 from **0.6512** $\rightarrow$ **0.6777** (**+2.65** macro-F1; same calibrated reference as **gnn07** / zeroshot in this sweep). CLIP embeddings appear **complementary** to frozen VLM scores on the bipartite/co-error heads even when logits are calibrated post hoc.

6.4 Effect of bipartite topology (M3 $\rightarrow$ M4)

Replacing the $C \times C$ homogeneous graph with **bipartite attribute $\rightarrow$ object** flow lifts calibrated test macro-F1 from **0.6777** $\rightarrow$ **0.6889** (**+1.12** vs gnn12 in `reports/comparison/overall.json`). The qualitative structural argument—a single imaging “object” and many latent findings—remains the same (`scripts/gnn_bipartite.py`), with `mode=all` default edge weights unless `vlm_positive` is tuned.

6.5 Per-protocol observations

- The MLP baseline closes most of the gap to the residual GNN under the calibrated protocol; the structural prior (label graph) does **not** add measurable value on top of class-bias correction *unless* image features are also re-injected.
- **gnn12** / **gnn13** gain substantially when thresholds are constrained to **calib**: compare default split test macro-F1 at reported per-class thresholds (same split as tuning in that column—**still optimistic**) **0.6013**, versus leakage-free calibrated test **0.6777** / **0.6889**.

- Zeroshot @0.5 remains ≈ 0.047 macro-F1 (test)—the pathology of uniformly high thresholds against negatively biased logits. With **honest calibrated** thresholds on frozen probabilities only, zeroshot reaches ≈ 0.6512 , and the best bipartite head still delivers $\approx +0.038$ macro-F1 on top; **trained** adapters achieve usable masked macro-F1 at @0.5 on the defaults split (§6.1 first columns).

7. Ablations and Analysis

7.1 Edge mode in the bipartite GNN (`--edge_mode {all, vlm_positive} × --vlm_tau`)

`vlm_positive` masks out attributes whose VLM probability is below `vlm_tau`, with an all-ones fallback for empty rows. In our runs `mode=all` was used as the default; `mode=vlm_positive` adds a per-row sparsification that is most useful when the VLM is well-calibrated for *some* classes (No Finding, Cardiomegaly) and noisy for others (Pneumonia). It is exposed in `scripts/13_train_bipartite_gnn_adapter.py` for dataset-specific tuning.

7.2 Number of bipartite layers

`--gnn_hidden_dims 512,256` (i.e., $L = 2$) is the default. The first layer encodes high-frequency attribute $\rightarrow$ object updates, the second consolidates them. Going deeper (e.g., `512,256,128`) does not help because the bipartite graph has diameter 2 and additional layers only re-mix the same object state.

7.3 Co-error graph `top_k` and τ

`top_k=3`, $\tau=0.02$ was used as the default. Lower τ introduces noisy edges (rare co-errors); higher `top_k` densifies the graph and effectively averages neighbours, which hurts the residual GNN’s precision on rare classes (Pneumonia, Consolidation). The `gnn07` configuration is *small*, so over-densifying the graph quickly approaches a global mean and erases per-class structure.

7.4 α (residual scale)

All three GNNs default to $\alpha=0.5$. Setting $\alpha=0$ reduces the model to a pure label-graph predictor with no VLM anchor and underperforms (`gnn07` cannot recover from the VLM’s negative bias without the residual). Setting $\alpha=1.0$ makes the adapter a strict additive correction and is more stable but slightly less accurate than 0.5 on our splits.

7.5 Best-checkpoint metric (`--best_metric`)

`val_bce` is the default. `val_macro_f1_05` selects on the same metric we report at @0.5 and biases the model toward overconfident logits. `val_macro_f1_thr` selects on per-class-tuned val F1 and is the strongest in-sample metric, which is exactly why we **do not** use it: it is one threshold-tuning step away from the same leakage as §6.2.

7.6 Class-weighted loss (pos_weight_max=100)

The 7 CheXpert classes are extremely imbalanced (Pneumonia $\approx$ 2% positive in our train split, No Finding $\approx$ 50%). Capping pos_weight at 100 prevents a single rare class from dominating the gradient. Removing the cap (pos_weight_max= ∞) destabilises early training and makes the cosine schedule diverge.

7.7 The flat-MLP control

The MLP baseline already receives the *full* [logits; probs] vector and spans per-class affine-ish recalibration. It scores **0.6544** on calibrated test (**reports/comparison/overall.js**) so calibrated macro-F1 gains **above that line** (> 0.6544) are evidence of *cross-class structure* (CLIP + graph) rather than scalar recalibration alone.

8. Discussion

8.0 Framing: when to prefer a GNN adapter over VLM fine-tuning

Use **adapter-only domain adaptation** (no VLM gradients) when: (i) you cannot afford full fine-tuning compute or liability of changing a shared foundation model; (ii) you need **fast iteration** on label policies, thresholds, or hospital-specific biases; (iii) you already have **offline VLM scores** and want a reproducible head. Prefer a **GNN head** over a flat MLP when you believe **label structure** (co-occurrence, co-error, or image-attribute factorization) should be **baked in** rather than re-learned from limited data—and when you can supply a graph prior (co-error from train) or a bipartite template (attribute→object). Prefer **full VLM fine-tuning** when the domain gap is primarily **visual** (e.g., completely new modality or resolution) and no auxiliary frozen image encoder closes that gap; this repo’s best results suggest that **re-injecting CLIP** captures part of that visual gap *without* touching the VLM.

8.1 What actually drives the gains

In ascending order of contribution to **0.6889** calibrated bipartite macro-F1 **vs** calibrated frozen probs (≈ 0.6512) and naïve **@0.5**:

1. **Class-weighted masked BCE plus honest threshold protocol** (vs naïve @0.5): lifts calibrated frozen probabilities from 0.047 $\rightarrow \approx 0.65$ macro-F1 (zeroshot; see §6.1 footnote on export_rows_to_predictions + calib-only thresholds) and aligns adapter scores with decision-rule hygiene.
2. **CLIP image features re-injected** (M2 $\rightarrow$ M3 vs the same ≈ 0.6512 calibrated reference): **+2.65** macro-F1 (to **0.6777**) under 08/09.
3. **Bipartite attribute $\rightarrow$ object topology** (M3 $\rightarrow$ M4): **+1.1** F1 by matching the data-generating process.
4. **Co-error label graph** (M1 $\rightarrow$ M2): essentially neutral once calibration is honest. The graph encodes a prior that the adapter can also learn from data given enough capacity.

8.2 What does *not* drive gains

- The choice of gnn07 adjacency normalization does not matter beyond row-stochastic.
- Going past 2 bipartite layers does not improve results.
- Selecting checkpoints on val-thr macro F1 inflates val numbers without a matching test improvement.

8.3 Practical recommendation

For a production system that must keep the foundation VLM frozen, **gnn13_clip_bipartite evaluated under the 4-way calibrated protocol** is the recommended configuration (**0.6889** calibrated test macro-F1 in reports/comparison/overall.json). If CLIP inference is unacceptable, **vlm_mlp under that same protocol remains competitive (0.6544)** at lower architecture complexity.

9. Limitations

1. **Single VLM source.** All non-zeroshot adapters consume Qwen2-VL-2B-Instruct outputs; we have not measured robustness under a different VLM.
 2. **Single CLIP backbone.** openai/clip-vit-base-patch32 is the smallest CLIP and is not pretrained on chest X-rays; a domain-adapted encoder (e.g., Biomed-CLIP) would likely lift M3/M4 further.
 3. **Single random seed.** All numbers are seed-42; we have not reported variance across seeds.
 4. **Patient-grouped but not site-grouped splits.** CheXpert is single-site, so distribution shift is not measured here.
 5. **Macro-F1 only.** AUROC, balanced accuracy and per-class recall at fixed precision would round out the comparison.
 6. **Threshold grid is coarse** (step=0.05); a finer grid or differentiable-threshold method (e.g., F β -soft) might shave a small additional F1.
-

10. Reproducibility Summary

- Code: scripts/01-13, scripts/common_multilabel.py, scripts/model_registry.py, scripts/gnn_bipartite.py.
- Configs: configs/{data,graph,train_gnn,train_clip_gnn,eval}.yaml.
- Seeds: 42 throughout (set_seed).
- Splits: data/processed/splits/ (3-way) and data/processed/splits_4way/ (4-way).
- Graph: data/processed/graph/{edge_index,edge_weight}.json.
- CLIP caches: data/processed/embeddings/clip_vitb32_default.pt and clip_vitb32_calibrated4way.pt (one per split regime; **never** mix row sets in a single cache file).
- Run registry: data/processed/experiments/<model_id>/<protocol>/{runs_index,latest,be

- Reports: reports/gnn_adapter/report.md, reports/comparison/overall.{md,json}.
- Inference UI: gradio_inference.py (resolves best.json under .../<model_id>/<default|calib optional legacy weights may live in .../_archive_20260503/ after cleanup).

11. Conclusion

We argue that **GNN-based adapters remain a practical recipe for domain adaptation in multi-label CXR classification without fine-tuning the VLM**: the foundation model supplies fixed (z, p) scores (CLIP optional), while a lightweight head learns **structured, label-aware corrections**. On the replicated registry bundle (**RUN_ID=repro_full_20260503**, reports/comparison/overall.json), bipartite **gnn13_clip_bipartite** attains **0.6889** calibrated test macro-F1— $\approx$ **+3.4** over calibrated **vlm_mlp (0.6544)** and $\approx$ **+3.8** over calibrated frozen probabilities alone (0.6512), while $t=0.5$ keeps masked macro-F1 near **0.05** on zeroshot without threshold optimisation. Threshold tuning **on evaluation data** still inflates **validation** residuals by $\approx +0.61$ macro-F1 (see *_LEAKY.json / §6.2); reserving **calib** removes that ambiguity. Jointly these facts support **“freeze VLM, adapt with graphs (and calibrated decisions)”** whenever full encoder fine-tuning is prohibitive.

Appendix A — Script-to-Artifact Map

Stage	Script	Primary artifacts
Canonical labels	01_build_canonical_labels.py	data/processed/multilabel/canonical_labels.py
VLM alignment	02_align_vlm_outputs.py	data/processed/multilabel/aligned_vlm_outputs.py
3-way splits	03_make_multilabel_splits.py	data/processed/splits/{train,val,test}_3way.py
4-way splits	03_make_multilabel_splits.py	data/processed/splits_4way/{train_val_test}_4way.py
Co-error graph	04_build_coerror_graph.py	data/processed/graph/ (defaults) plus data/processed/graph_4way/ (4-way splits)
Frozen VLM eval	05_run_baseline_frozen_vlm.py	baseline_frozen_vlm/metrics.json
MLP baseline	06_run_baseline_mlp.py	.../vlm_mlp/<protocol>/<run_id>/{metrics.json, ...}
Residual GNN	07_train_gnn_adapter.py	.../gnn07_label_residual/<protocol>/<run_id>/{metrics.json, ...}
Threshold tuning	08_tune_thresholds.py	.../<protocol>/<run_id>/per_class_thresholds.json (calib-fed) or *_tuned_on_val_LEAKY.json
Evaluation	09_evaluate_test.py	.../<run_id>/test_metrics_calibrated.json (honest) or *_LEAKY.json (diagnostic)
Ablation collation	10_run_ablations.py	data/processed/experiments/ablations/ (writes new dir if missing)
Markdown report	11_package_report.py	reports/{gnn_adapter/report.md, comparison/overall.{md,json}}

Stage	Script	Primary artifacts
CLIP+VLM GNN	12_train_clip_vlm_gnn_adapter/gnn12_clip_vlm_homo/<protocol>/	
Bipartite GNN	13_train_bipartite_gnn_adapter/gnn13_clip_bipartite/<protocol>/	

Appendix B — Default Hyperparameters

Model	Hidden	Layers	LR	Sched	Epochs	Batch	Dropout α		Other
vlm_mlp	64	2	1e-3	none	20	full	0.1	—	AdamW, wd 1e-4, pos_weight≤100
gnn07_label_resolution	128	1 dualg-pass	3e-4	cosine + warmup 2	≤80 (es 18)	full	0	0.5	grad_clip 1.0, val_bce checkpoint
gnn12_clip_vlm_homo GNN	64	16	3e-4	cosine + warmup 2	≤60 (es 16)	32	0	0.5	CLIP B/32 frozen, cached embeddings
gnn13_clip_bipartite	256	16 bipartite	3e-4	cosine + warmup 2	≤60 (es 16)	32	0.2	0.5	object_dim 512, edge_mode=all

Appendix C — References

Grouped by topic for quick lookup; alphabetical within each topic.

C.1 Dataset & task

- Irvin J., Rajpurkar P., Koh M., Yu Y., Cicurel S., Chute C., ... Ng A. Y. (2019). *CheXpert: A Large Chest Radiograph Dataset with Uncertainty Labels and Expert Comparison*. Proceedings of AAAI Conference on Artificial Intelligence (AAAI).

C.2 Vision-language foundation models & parameter-efficient tuning

- Radford A., Kim J. W., Hallacy C., Ramesh A., Goh G., Agarwal S., ... Sutskever I. (2021). *Learning Transferable Visual Models From Natural Language Supervision*. Proceedings of ICML.

- Wang P., Bai S., Tan S., Wang S., Fan Z., Bai J., ... Zhou J. (Qwen team) (2024). *Qwen2-VL: Enhancing Vision-Language Model's Perception of the World at Any Resolution*. arXiv:2409.12191 <https://arxiv.org/abs/2409.12191>.
- Hu E. J., Shen Y., Wallis P., Allen-Zhu Z., Li Y., Wang S., Wang L., Chen W. (2022). *LoRA: Low-Rank Adaptation of Large Language Models*. Proceedings of ICLR.
- Houlsby N., Giurui A., Jastrzebski S., Morrison Q., Larochelle H., Gesmundo M., Attariyan H., Gelly S. (2019). *Parameter-Efficient Transfer Learning with NLP Adapter Modules*. Proceedings of ICML.

C.3 Graph neural networks — homogeneous convolution & message passing

- Kipf T. N., Welling M. (2017). *Semi-Supervised Classification with Graph Convolutional Networks*. ICLR poster.
- Hamilton W., Ying Z., Leskovec J. (2017). *Inductive Representation Learning on Large Graphs*. NeurIPS.
- Veličković P., Cucurull G., Casanova A., Romero A., Liò P., Bengio Y. (2018). *Graph Attention Networks*. ICLR.
- Gilmer J., Schoenholz S. S., Riley P., Vinyals O., Dahl G. E. (2017). *Neural Message Passing for Quantum Chemistry*. ICML. (Defines a general pairwise message/update framework that subsumes many bipartite and heterogeneous aggregations.)

C.4 Heterogeneous, relational & bipartite-style graphs

- Schlichtkrull M., Kipf T. N., Bloem P., van den Berg R., Titov I., Welling M. (2018). *Modeling Relational Data with Graph Convolutional Networks*. European Semantic Web Conference (ESWC). https://doi.org/10.1007/978-3-319-93417-4_38
- Wang X., Ji H., Shi C., Wang B., Ye Y., Cui P., Yu P. S. (2019). *Heterogeneous Graph Attention Network*. The Web Conference (WWW). <https://doi.org/10.1145/3308558.3313562>

C.5 Multi-label image recognition via label graphs

- Chen Z.-M., Wei X.-S., Wang P., Guo Y. (2019). *Multi-Label Image Recognition with Graph Convolutional Networks*. Proceedings of IEEE/CVF CVPR.

C.6 Threshold choice, probabilistic calibration & “leakproof” evaluation protocol

Choosing per-class thresholds and reporting accuracy/F1 mixes **classification calibration** with **selection bias** whenever the same labeled split is reused for tuning and leaderboard reporting.

- Lipton Z. C., Elkan C., Narayanaswamy B. (2014). *Optimal Thresholding of Classifiers to Maximize F1 Measure*. In *Proceedings of ECML PKDD 2014*

(Springer LNCS vol. 8726), pp. 225–239. https://doi.org/10.1007/978-3-662-44851-9_15 (related preprint: *Thresholding Classifiers to Maximize F1 Score*, arXiv:1402.1892.)

- Lewis D. D. (1995). *Evaluating and optimizing autonomous text classification systems*. ACM SIGIR. (Classic framing of precision/recall trade-offs via thresholds.)
- Platt J. C. (1999). *Probabilistic outputs for Support Vector Machines and comparisons to regularized likelihood methods*. In *Advances in Large Margin Classifiers* (MIT Press). (Platt / temperature-style scaling lineage.)
- Niculescu-Mizil A., Caruana R. (2005). *Predicting Good Probabilities with Supervised Learning*. ICML.
- Guo C., Pleiss G., Sun Y., Weinberger K. Q. (2017). *On Calibration of Modern Neural Networks*. Proceedings of ICML.
- Varma S., Simon R. (2006). *Bias in Error Estimation When Using Cross-Validation for Model Selection*. BMC Bioinformatics, 7(1), 91. <https://doi.org/10.1186/1471-2105-7-91>
- Cawley G. C., Talbot N. L. C. (2010). *On Over-fitting in Model Selection and Subsequent Selection Bias in Performance Evaluation*. Journal of Machine Learning Research (JMLR), 11, 2079–2107. <http://jmlr.org/papers/v11/cawley10a.html>
- Dietterich T. G. (1998). *Approximate Statistical Tests for Comparing Supervised Classification Learning Algorithms*. Neural Computation, 10(7), 1895–1923. (Multiple runs / paired tests when comparing adapters.)

Together, Varma & Simon (2006) and Cawley & Talbot (2010) underpin the methodological point that **any post hoc rule fit on labeled data—including per-class thresholds—must consume a disjoint hold-out** if the same numerical split is otherwise used for model comparison; our **calib** split instantiates exactly that separation from `train_fit`, `val`, and `test`.